




CSE308 Operating Systems

Concurrency & Synchronization

Dr S. Rajarajan

SASTRA

Concurrent Processes



Types of processes

- An **independent process** is one that executes independently without interacting or affecting other processes
- A **cooperating process** is one that can affect or be affected by other processes executing in the system
- Cooperating processes may get into conflicts

- Cooperating processes can
 - either **directly share a logical address space** that is, both code and data (**threads**)
 - or be allowed to **share data** only through files
 - or communicate through messages. (**pipes, shared memory, message queues**)
- In all the cases, concurrent access to **shared data** may result in **data inconsistency**.

- Shared resources
 - Files
 - IPC mediums – shared memory & message queue
 - Memory buffers
 - Variables
 - Hardware device e.g printer
 - Database

Concurrent Vs Parallel processing

- **Concurrent processing**
 - Under a uniprocessor, the scheduler keeps the CPU to switch from process to process to avoid the idle time of CPU.
- **Parallel processing**
 - Under a multi-core / multi-processor system either two separate processes or threads of a same process can be parallelly executed on different **cores or processors** at a time.

- **concurrent or parallel execution** can contribute to issues involving the **integrity of data** shared by several processes.
- We will discuss **various mechanisms** to ensure the **orderly execution** of cooperating processes so that **data consistency** is maintained.


Bounded buffer producer-consumer problem

- **Producer** is a process which **produces data** and **puts them into a buffer**.
- **Consumer** is a process that **retrieves** and **consumes data** from the buffer.
- If **buffer is full** then **producer must not add data**.
- If **buffer is empty** then **consumer must not attempt to retrieve data**.

Producer

- Counter tells number of data in buffer & it is initialized to 0

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE); //Buffer full  
        /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```



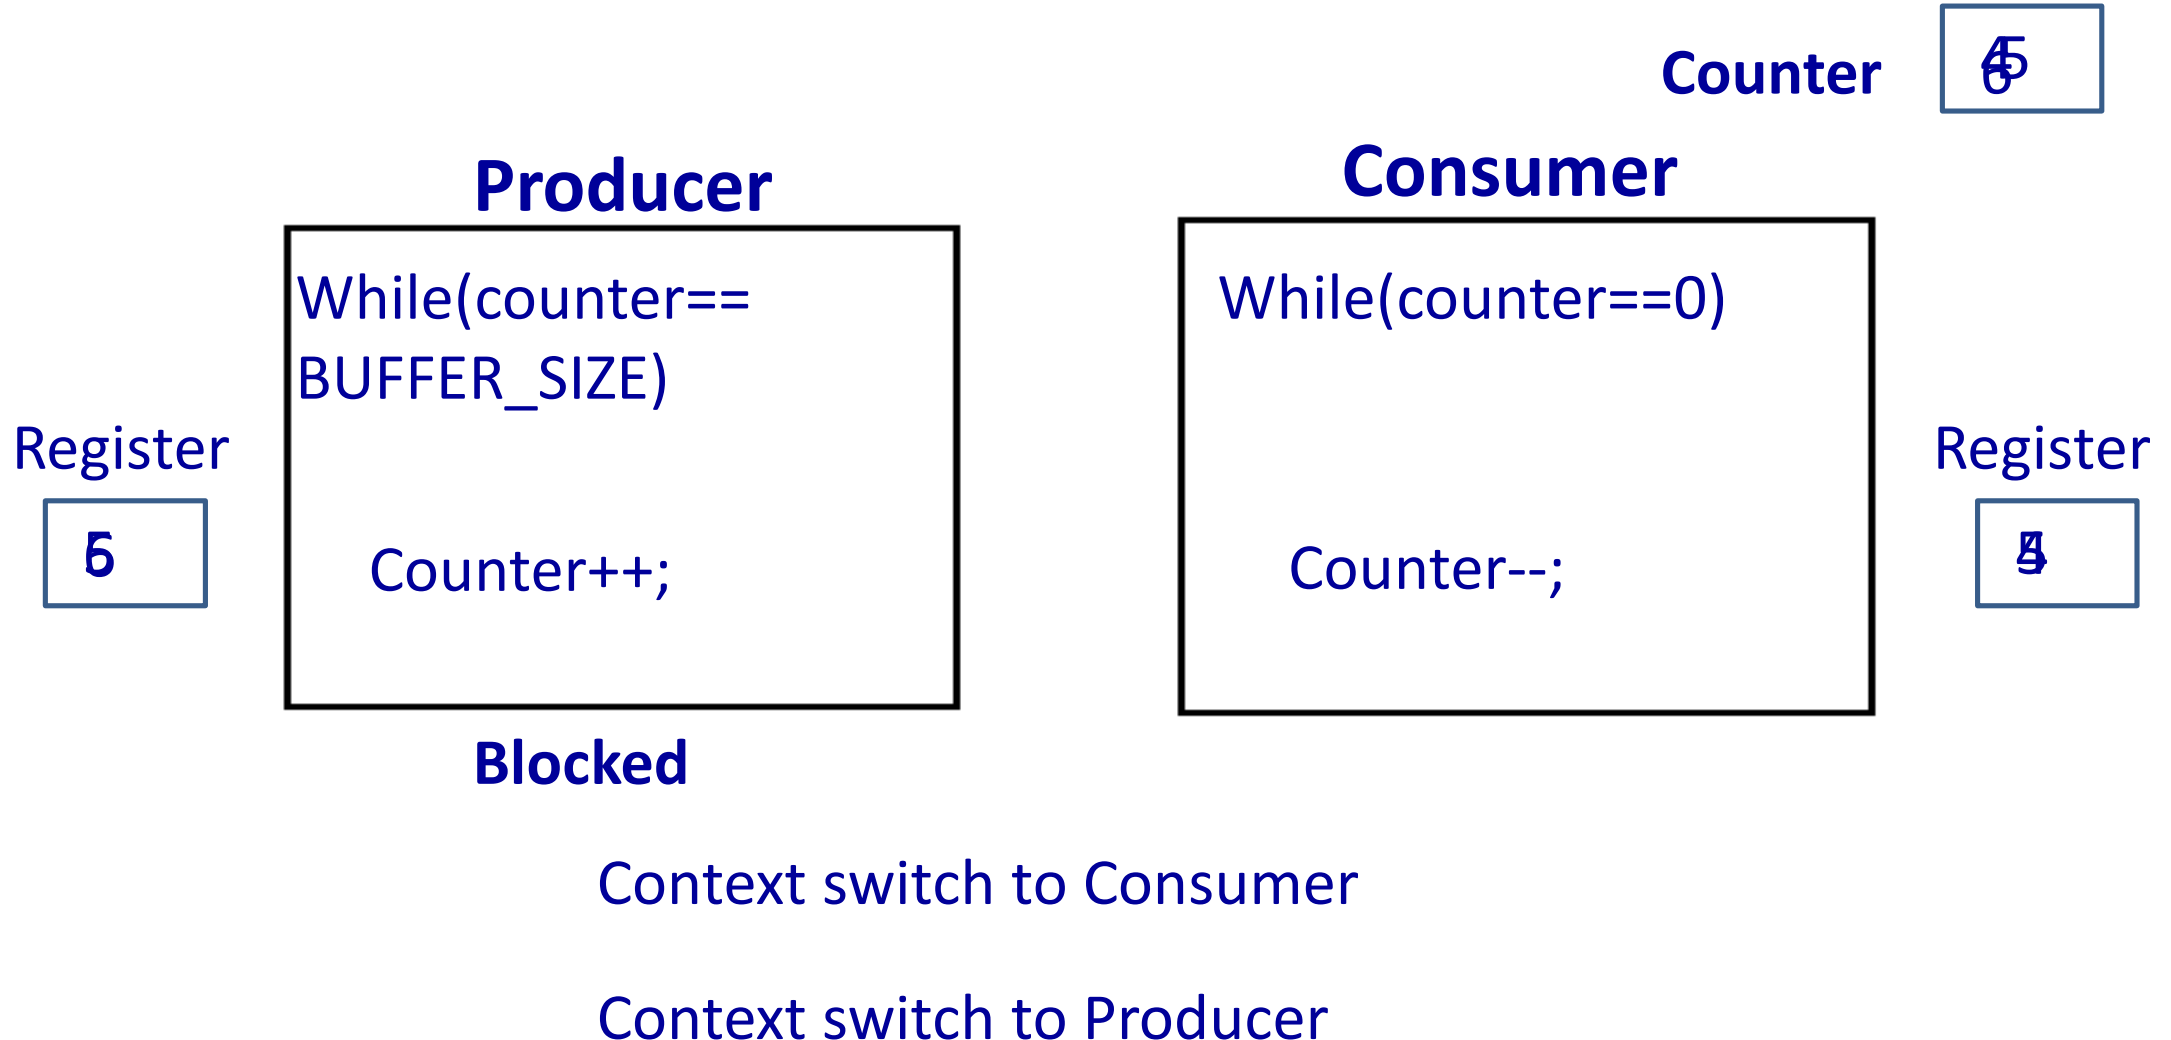
If in is 6
 $(6+1) \% 7 = 0$

Consumer

```
while (true) {  
    while (counter == 0) //Buffer empty  
        ; /* do nothing */  
  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter-- ;  
  
    /* consume the item in next_consumed */  
}
```

- The producer and consumer routines are **correct** if they are **executed sequentially**
- **But** they may function **incorrectly** when executed **concurrently**
- Suppose that the value of the variable **counter** is **5** and that the producer and consumer processes **concurrently execute** the statements “**counter++**” and “**counter--**”.
- The **correct result** is **counter = 5**, which is generated correctly if the producer and consumer execute separately
- But under concurrent execution of the two statements, the value of the variable counter may be **4, 5, or 6**.

We can show that the value of counter may become incorrect as follows



- Note that the statement “counter++” may be implemented in machine language (on a typical machine) as follows:
- *register1 = counter*
- *register1 = register1 + 1*
- *counter = register1*

- *Similarly, the statement “counter--” is implemented as follows:*
- *register2 = counter*
- *register2 = register2 – 1*
- *counter = register2*
- Even if register1 and register2 are the same physical register (an accumulator, say), remember that the **contents of this register** will be **saved and restored** by context switching.

- The concurrent execution of “counter++” and “counter--” is equivalent to a sequential execution in which the lower-level statements presented previously are interleaved in some arbitrary order.
- One such interleaving is the following:

T_0 :	<i>producer</i>	execute	$register_1 = counter$	$\{register_1 = 5\}$
T_1 :	<i>producer</i>	execute	$register_1 = register_1 + 1$	$\{register_1 = 6\}$
T_2 :	<i>consumer</i>	execute	$register_2 = counter$	$\{register_2 = 5\}$
T_3 :	<i>consumer</i>	execute	$register_2 = register_2 - 1$	$\{register_2 = 4\}$
T_4 :	<i>producer</i>	execute	$counter = register_1$	$\{counter = 6\}$
T_5 :	<i>consumer</i>	execute	$counter = register_2$	$\{counter = 4\}$

- Notice that we have arrived at the incorrect state “**counter == 4**”, indicating that four buffers are full, when, in fact, **five buffers** are full.
- If we reversed the order of the statements at *T4 and T5*, *we would arrive at the incorrect state “**counter == 6**”*.
- We would arrive at this incorrect state because we allowed both processes to **manipulate** the variable **counter concurrently**.

Multiprocessor System

Producer

```
While(counter==  
BUFFER_SIZE)
```

```
Counter++;
```

Register

6

Consumer

```
While(counter==0)
```

```
Counter--;
```

Register

4

Counter

5

Race condition

- A **race condition** or **race hazard** is the condition of an software system where the system's substantive behaviour is dependent on the sequence or timing of other **uncontrolled events**, leading to **unexpected or inconsistent results**.
- Situations where several processes access and **manipulate** the same **data concurrently** and the **outcome** of the execution depends on the **particular order** in which the access takes place leads to a **race condition**.
- To guard against the race condition, we need to ensure that **only one process** at a time can be manipulating the variable counter.
- To make such a guarantee, we require that the **processes be synchronized** in some way

Sequential access

Thread 1	Thread 2		Integer value
			0
read value		←	0
increase value			0
write back		→	1
	read value	←	1
	increase value		1
	write back	→	2

Concurrent / Parallel access

Thread 1	Thread 2		Integer value
			0
read value		←	0
	read value	←	0
increase value			0
	increase value		0
write back		→	1
	write back	→	1

What is synchronization ?

- Situations such as the one just described occur frequently in operating systems as different parts of the system **manipulate same resources**.
- The growing importance of **multi-core systems** has brought an increased emphasis on developing **multithreaded** applications.
- In such applications, **several threads** —which are quite possibly **sharing data** —are running in **parallel** on different processing cores
- When two processes want to access a **shared resource** there should be certain **discipline, monitoring** and **regulation** imposed while accessing that resource.

Starting point - Resource sharing



Synchronization – enforcing order on resource usage



When no Synchronization enforced



Consequence - conflicts



Limits on number of simultaneous users for a resource



Race condition - example

- Suppose that two processes, **P1** and **P2**, share the **global variable a**.
- At some point in its execution, **P1** updates **a=1**, and at some point in its execution, **P2** updates **a= 2**.
- Thus, the two tasks are **in a race** to write variable
- In this example the “**looser**” of the race (the process that **updates last**) determines the **final value of a**.

Example 1

int a

Process 1

a=1

Process 2

a=2

- If the sequence of execution of the two processes is **P1 followed by P2** then 'a' will be modified as
 - P1: a=1 & P2: a=2
 - So the final value of 'a' will be 2
- If the sequence of execution of the two processes is **P1 followed by P2** then 'a' will be modified as
 - P2: a=2 & P1: a=1
 - Then the final value of 'a' will be 1

Example 2

Process P3
 $b = b + c$

$b=1 \quad c=2$

Process P4
 $c = b + c$

- Though the two processes update different variables still the final values of the b & c depends on the order in which two processes are executed.
- If the order is **P3 followed by P4** then $b=3, c=5$
- If the order is **P4 followed by P3** then $b=4, c=3$

- Count is a global variable initialized with '0'

P1

```
If( Count == 0)
{
    .....
    Count++;
}
Count=0
```

P2

```
If( Count == 0)
{
    .....
    Count++;
}
Count=0
```

- Under **sequential execution**, only one process could increment Count and so it will become 1.
- But under **concurrent or parallel processing** count might become **2**

P1

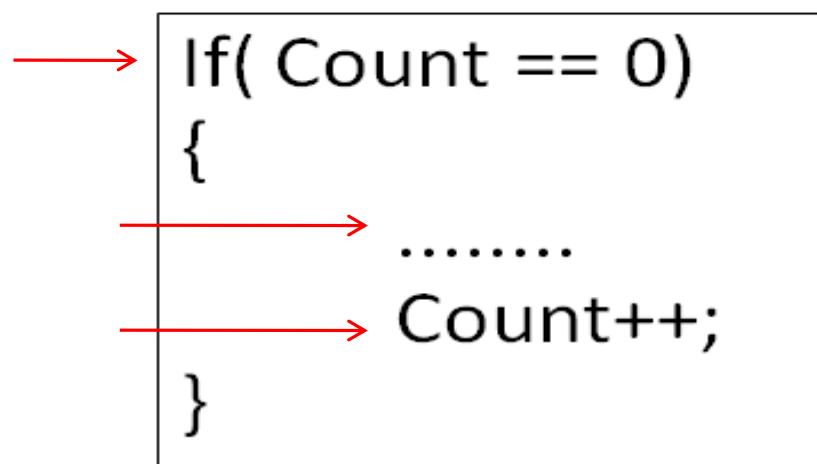


Diagram illustrating the execution flow of process P1. The code is enclosed in a box. Red arrows indicate the sequence of execution: one arrow points to the opening curly brace '{', another to the condition 'If(Count == 0)', a third to the ellipsis '.....', and a fourth to the increment statement 'Count++;'. A fifth arrow points to the closing curly brace '}'.

```
If( Count == 0)
{
.....
Count++;
}
```

P2

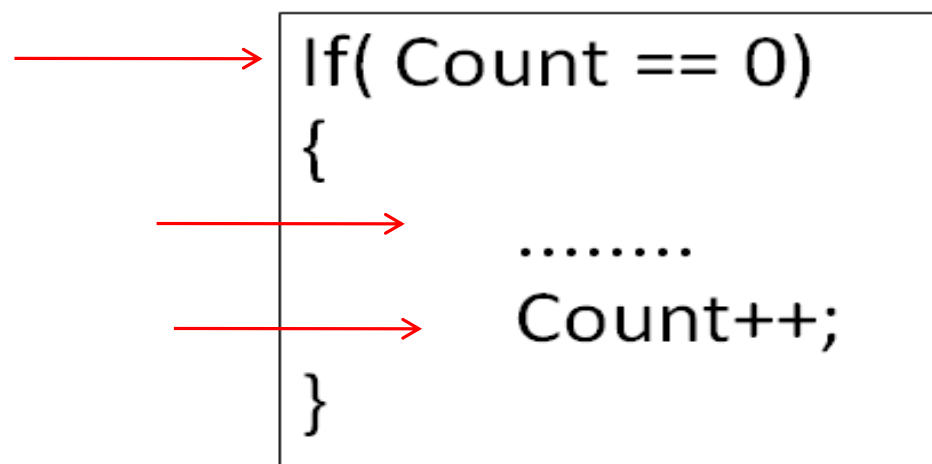


Diagram illustrating the execution flow of process P2. The code is enclosed in a box. Red arrows indicate the sequence of execution: one arrow points to the opening curly brace '{', another to the condition 'If(Count == 0)', a third to the ellipsis '.....', and a fourth to the increment statement 'Count++;'. A fifth arrow points to the closing curly brace '}'.

```
If( Count == 0)
{
.....
Count++;
}
```

The Critical-Section

- Consider a system consisting of n processes $\{P_0, P_1, \dots, P_{n-1}\}$.
- Each process has a **segment of code**, called a **critical section**, in which the process may be **changing common variables**, updating a table, writing a file, and so on.
- The important feature of the system is that, when one process is **executing in its critical section**, no other process is allowed to execute in its critical section.
- **That is, no two processes can be executing in their critical sections at the same time.**



Process

remainder_section

critical section

entry_section

access shared resource

exit_section

remainder_section


```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE);  
        /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    CS counter++;  
}
```

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    CS counter-- ;  
  
    /* consume the item in next_consumed */  
}
```

Requirements

- The ***critical-section problem is to design a*** protocol that coordinates the cooperative processes.
- Each process **must request permission** to enter its critical section.
- **Entry section-** The section of code that begins to access the shared resource
- **Exit section-** Section of code following the shared resource usage
- The remaining code is the **remainder section**

E.g. Critical section for resource R1



Strategy

- ***Critical Resource:-*** A resource that may lead to conflict between processes.
 - It can be a variable or file or hardware device or memory region
- **Critical Section:-** The portion of the program/process that accesses the critical resource
- **Mutual exclusion:-** Only one process at a time to be allowed to enter into its critical section

Critical section conditions

- A solution to the critical-section problem must satisfy the following three requirements:
- **1. Mutual exclusion.** If process *P* is executing in its critical section, then no other processes can be executing in their critical sections.
- **2. Progress.** If no process is in its critical section then an intending process should be able to enter provided it is the only process entering. This selection cannot be postponed indefinitely.
- **3. Bounded waiting.** There exists a bound, or limit, on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted (***no starvation***)

Kernel level race conditions

- At a given point in time, many kernel-mode processes may be active in the operating system.
- As a result, the code implementing an operating system (**kernel code**) is subject to several possible race conditions.
- E.g file management – two kernel modules updating same file
- Two general approaches are used in operating systems:
preemptive kernels and **non-preemptive kernels**.
- A **preemptive kernel** allows a process to be preempted while it is running in kernel mode.
- A **non-preemptive kernel** does not allow a process running in kernel mode to be preempted; a kernel-mode process will run until it exits kernel mode, blocks, or voluntarily yields control of the CPU.

- Obviously, a **non-preemptive kernel** is essentially **free from race conditions** on kernel data structures, as **only one process is active** in the kernel **at a time** (no context switching among concurrent processes).
- We cannot say the same about **preemptive kernels**, so they must be **carefully designed** to ensure that shared kernel data are free from race conditions.
- **Preemptive** kernels are especially **difficult** to design for **SMP architectures**, since in these environments it is possible for **two kernel-mode processes** to run simultaneously on **different processors**

1. Peterson's Solution

- A classic **software-based** solution to the critical-section.
- Because of the way **modern computer architectures** perform machine-language instructions- **load and store**, there are **no guarantees** that Peterson's solution will work correctly on such architectures.
- However, it provides a **good algorithmic description** of solving the **critical-section problem** and illustrates some of the **complexities involved** in designing software that addresses the requirements of **mutual exclusion, progress, and bounded waiting**

- Peterson's solution is **restricted to two processes** that alternate execution between their critical sections and remainder sections.
- The processes are numbered **P_0 and P_1** .
- *For convenience, when presenting P_i , we use P_j to denote the other process; that is, j equals $1 - i$.*
- Peterson's solution requires the two processes to share two data items:
 - **int turn;**
 - **boolean flag[2];**
- The variable **turn** indicates **whose turn** it is to enter its critical section.
- That is, if **turn == i**, then process **P_i is allowed to execute in its critical section.**

- The **flag array** is used to **indicate** if a **process wants** to enter its **critical section**.
- For example, if **flag[i] is true**, this value indicates that ***Pi is willing to enter*** its critical section.
- To enter the critical section, P_i ,
 - *first sets flag[i] to be true*
 - *and then sets turn to the value j*, thereby asserting that if the other process wishes to enter the critical section, it can do so.
- If **both processes try** to enter at the same time, turn will be set to **both i and j** at roughly the same time.
- Only one of these assignments will last; the other will occur but will be **overwritten immediately**.
- The **eventual value of turn determines** which of the two processes is allowed to enter its critical section first.

Peterson's algorithm

do {

```
flag[i] = true;  
turn = j;  
while (flag[j] && turn == j) ;
```

If condition is false
then process enters
the critical section

→ critical section

```
flag[i] = false;
```

remainder section

} while (true);

Both P0 & P1 wants to enter CS

Flag

0	1
false	false

turn

0

```
P0:    flag[0] = true;
P0_gate: turn = 1;
        while (flag[1] == true && turn == 1); False
            // busy wait
        // critical section
        ...
        // end of critical section
        flag[0] = false;
```

```
P1:    flag[1] = true;
P1_gate: turn = 0;
        while (flag[0] == true && turn == 0); False
            // busy wait
        // critical section
        ...
        // end of critical section
        flag[1] = false;
```

P1 in CS & P0 wants to enter

Flag 0 1
 false true
turn 0

```
P0:    flag[0] = true;
P0_gate: turn = 1;
        while (flag[1] == true && turn == 1); True
                // busy wait

        // critical section
...
// end of critical section
flag[0] = false;
```

```
P1:    flag[1] = true;
P1_gate: turn = 0;
        while (flag[0] == true && turn == 0);
                // busy wait
        → // critical section
...
// end of critical section
flag[1] = false;
```

Only P0 wants to enter

	0	1
Flag	false	false
turn	0	

```
P0:    flag[0] = true;
P0_gate: turn = 1;
        while (flag[1] == true && turn == 1); False
                                // busy wait
→ // critical section
...
// end of critical section
flag[0] = false;
```

```
P1:    flag[1] = true;
P1_gate: turn = 0;
        while (flag[0] == true && turn == 0);
                                // busy wait
        // critical section
...
// end of critical section
flag[1] = false;
```

Wrong turn

Both P0 & P1 wants to enter CS

Flag

0	1
false	false

turn

0

```
P0:    flag[0] = true;
P0_gate: turn = 0;
        while (flag[1] == true && turn == 1); False
            // busy wait
→ // critical section
...
// end of critical section
flag[0] = false;
```

```
P1:    flag[1] = true;
P1_gate: turn = 1;
        while (flag[0] == true && turn == 0); False
            // busy wait
→ // critical section
...
// end of critical section
flag[1] = false;
```


Wrong turns

Both P0 & P1 wants to enter CS

	0	1
Flag	false	false
turn	0	

```
P0:    flag[0] = true;
P0_gate: turn = 0;
        while (flag[1] == true && turn == 0); True
                                // busy wait

        // critical section
...
// end of critical section
flag[0] = false;
```

```
P1:    flag[1] = true;
P1_gate: turn = 1;
        while (flag[0] == true && turn == 1); True
                                // busy wait

        // critical section
...
// end of critical section
flag[1] = false;
```

1. Situation 1: Currently turn is 0 , P1 has no interest in CS & P0 wants to enter CS ?

That means P0 changes Flag[0]=true, changes turn=1 and runs while (Flag[1] && turn =1). But since P1 has no interest, its Flag[1] is false. So P0 breaks while loop and enters CS

2. Situation 2: Currently turn is 0 , P1 has entered CS & P0 wants to enter CS ?

That means P0 changes Flag[0]=true, changes turn=1 and runs while (Flag[1] && turn =1). But since P1 has interest, its Flag[1] is true. Which means P1 is already in its CS. So P0 stuck in while loop .

3. Situation 3: Currently turn is 0 , P1 has interest in CS and has changed Flag[1]= true and it about to change turn=1. Meanwhile P0 changes Flag[0]=true and it is about to change turn=1. Which will enter CS ?

This is a typical race condition. Depending upon who changes turn at last, its value will be either 0 or 1. If turn becomes 0 then P0 enters CS else P1 enters.

- We now prove that this solution is correct. We need to show that:
 - 1. Mutual exclusion is preserved.
 - 2. The progress requirement is satisfied.
 - 3. The bounded-waiting requirement is met.

Mutual exclusion proof

- Both P0 and P1 can never be in the critical section at the same time
- If both processes are in their critical sections then the state must satisfy **flag[0]** and **flag[1]** are **true** and **turn = 0** and **turn = 1**.
- **No state** can satisfy both **turn = 0** and **turn = 1**, so there can be no state where both processes are in their critical sections.

Progress requirement proof

- To prove properties 2 and 3, we note that a process *P_i* **has to be prevented from** entering the critical section then it is only if gets stuck in the while loop with the condition **flag[j] == true and turn == j**; this loop is the only one possibility.
- If *P_j* **is not ready** to enter the critical section, then **flag[j] == false**, and *P_i* **can enter its** critical section.
- If *P_j* **has set flag[j] to true** and is also executing in its while statement, then either **turn == i**.
- The *P_i* changes **turn to j** and *P_j* **enters CS**

- If P_j enter critical section then when it *exits its critical section*, it will ***reset $flag[j]$ to false and also set turn to i*** allowing P_i to enter its critical section.
- So progress is guaranteed.

Bounded waiting proof

- Bounded waiting means that **the number of times a process is bypassed** by another process after it has indicated its desire to enter the critical section is bounded by a function of the number of processes in the system.
- In Peterson's algorithm, a process will **never wait longer than one turn** for entrance to the critical section.

2. Synchronization Hardware

- Software-based solutions such as Peterson's are **not guaranteed** to work on **modern computer architectures**.
- we explore several more solutions to the critical-section problem using techniques ranging from **hardware to software-based APIs**.
- All these solutions are based on the premise of **locking** — that is, protecting critical regions through the use of locks.

2.1 Interrupt disabling

- Critical-section problem could be solved easily in a **single-processor** environment if we could **prevent interrupts** from occurring while a **shared variable** was being modified – *non-preemptive*.
- In this way, we could be sure that the current sequence of instructions would be allowed to **execute in order** without **preemption** and **context switching**.
- This is often the approach taken by **non-preemptive kernels**.
- Unfortunately, this solution is not as feasible in a **multiprocessor** environment

Pseudo-Code

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```

2.2 Special hardware instructions

- Many modern computer systems provide special hardware instructions that allow us either
 - to test and modify the content of a word
 - or to swap the contents of two words **atomically**—that is, as one **uninterruptible** unit.
- We can use these special instructions to solve the critical-section problem in a relatively simple manner
- E.g. **test and set()** and **compare and swap()**
- instructions.

Compare & Swap Instruction

Checks a local value ***word** against a **testval**, if they are same, replaces with a newval

```
int compare_and_swap(int *word, int testval, int newval)
{
    int oldval;
    oldval = *word;

    if (oldval == testval) *word = newval;

    return oldval;
}
```

```
void p() { while( (c_a_s(bolt, 0, 1) == 1) ; CS }
```

The process that finds Bolt as 0 enters into the critical section

```

→ int Compare_and_Swap(int *word, int testval, int newval)
{
    int oldval;
    oldval = *word;
    if ( oldval == testval ) *word = newval;
    return oldval;
}

const int n= /* number of processes */
int bolt;
void P( In i )
{
    while( Compare_and_Swap(&bolt, 0, 1) != 1);
    /* do nothing */
    /* Critical Section */
    bolt = 0 ;
    /* remaining code */
}

void main()
{
    bolt = 0 ;
    parbegin(P(1), P(2), ...P(n));
}

```

bolt




```

int Compare_and_Swap(int *word, int testval, int newval)
{
    int oldval;
    oldval = *word;
    if ( oldval == testval ) *word = newval;
    return oldval;
}

const int n= /* number of processes */
int bolt;
void P( i )
{
    while( Compare_and_Swap( bolt, 0, 1 ) != 1 );
    /* do nothing */
    /* Critical Section */
    bolt = 0 ;
    /* remaining code */
}

void main()
{
    bolt = 0;
    parbegin(P(1), P(2), ...P(n));
}

```

bolt

1



test and set()

- The important characteristic of this instruction is that it is **executed atomically**.
- Thus, if two test and set() instructions are executed simultaneously (each on a different CPU), they will be **executed sequentially** in some **arbitrary order**.
- If the machine supports the test and set() instruction, then we can implement **mutual exclusion** by declaring a **boolean variable lock, initialized to false**.

Lock

false

```
boolean test and set(boolean *target) {
```

```
    boolean rv = *target; false
```

```
    *target = true;
```

```
    return rv; false
```

```
}
```

```
do { while ( test and set(&lock )); /* do nothing */
```

```
/* critical section */
```

```
lock = false;
```

```
/* remainder section */
```

```
} while (true);
```

- Although these algorithms satisfy the mutual-exclusion requirement, they **do not satisfy the bounded-waiting** requirement.
- We present another algorithm using the test and set() instruction that satisfies all the critical-section requirements.

```
do {  
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test and set(&lock); false  
    waiting[i] = false;  
    /* critical section */  
    j = (i + 1) % n; / To give the chance to next waiting process  
    while ((j != i) && !waiting[j]) / J is not waiting for CS  
        j = (j + 1) % n; / Continue search  
    if (j == i) / Unable to find a process  
        lock = false;  
    else // Next process 'j' to enter CS found  
        waiting[j] = false;  
} while (true);
```

Mutual exclusion proof

- Process *P_i* can enter its critical section only if either **waiting[i] == false** or **key == false**.
- The value of **key** can become **false** only if the **test and set()** is executed.
- The **first process** to execute the **test and set()** will find **key == false**; all others must wait.
- The variable **waiting[i]** can become **false** only if another process leaves its critical section; **only one waiting[i]** is **set to false**, maintaining the mutual-exclusion requirement

Progress proof

- The arguments presented for mutual exclusion also apply here.
- Since a process exiting the critical section either **sets lock to false** or **sets waiting[j] to false**.
- Both allow a process that is waiting to enter its critical section to proceed.

Bounded-waiting proof

- When a process leaves its critical section, it scans the array waiting in the cyclic ordering $(i + 1, i + 2, \dots, n - 1, 0, \dots, i - 1)$.
- *It designates the **first process** in this ordering that is in the entry section (**waiting[j] == true**) as **the next one to enter** the critical section.*
- Any process waiting to enter its critical section will thus do so within $n - 1$ turns.

3. Mutex Locks

- Hardware-based solutions are **complicated** as well as generally **inaccessible** to application **programmers**.
- Instead, ***operating-systems*** designers build software tools to solve the **critical-section problem**
- The simplest of these tools is the **mutex lock** (the term ***mutex*** is short for ***mutual exclusion***).
- That is, a **process must acquire the lock** before entering a critical section; it **releases the lock** when it exits the critical section.

- The **acquire()** function acquires the lock, and the **release()** function releases the lock.
- mutex lock has a **boolean variable** available whose value indicates if the **lock is available or not**.
- If the **lock is available**, a call to **acquire()** succeeds, and the lock is then considered unavailable.
- A process that attempts to acquire an **unavailable** lock is **made to wait until the lock** is released

- The definition of acquire() is as follows:

available = true

acquire()

{

while (!available); /* busy wait */

 available = false;

}

release()

{

 available = true;

}

available **false**

Applying mutex lock for CS

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

- Calls to either `acquire()` or `release()` must be performed **atomically**.
- Thus, mutex locks are often implemented using one of the **hardware mechanisms** described

Drawback - Busy waiting

- The main disadvantage of the implementation given here is that it requires **busy waiting**.
- While a process is in its critical section, any other process that tries to enter its critical section **must loop continuously** in the call to acquire().
- In fact, this type of mutex lock is also called a **spinlock because the process “spins”** while waiting for the lock to become available.
- We see the **same issue with** the code examples illustrating the **test and set()** instruction and the **compare and swap()** instruction.
- Busy waiting **wastes CPU cycles** that some other process might be able to use productively.

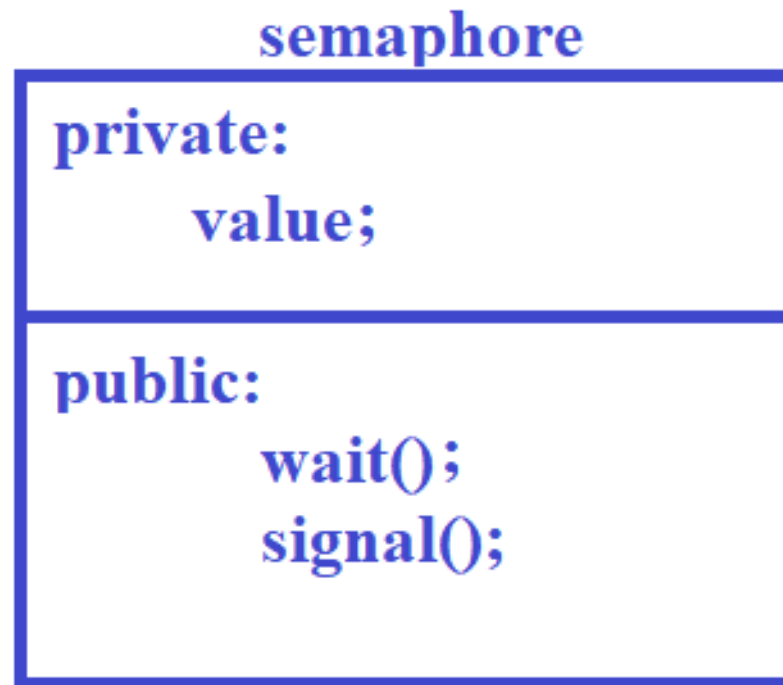
- **Spinlocks** do have an **advantage**.
- **No context switch is required** when a process must wait on a lock, and a context switch may take considerable time.
- Thus, when **locks are expected** to be held for **short times**, spinlocks are useful.
- They are often employed on **multiprocessor systems**.
- One thread can “**spin**” on one processor while another thread performs its critical section on another processor.

4. SEMAPHORE

Semaphores

- A more robust tool that can behave similarly to a mutex lock but **also provide more sophisticated ways** for processes to synchronize their activities.
- A **semaphore S** is an **integer** variable
 - that can be **initialized, incremented and decremented**.
 - and is accessed only through two standard **atomic** operations: **wait() (semWait)** and **signal() (semSignal)**.
- All modifications to the integer value of the semaphore in the **wait() and signal()** operations must be **executed indivisibly**.
- That is, when one process modifies the semaphore value, **no** other process can **simultaneously modify** that same semaphore value

- Semaphore can be treated **like an object** with s.value as its private data member and wait and signal as its public methods



- The definition of wait() is as follows:
- S=1

```
wait(S) {  
    while (S <= 0) ;// busy wait  
    S--;  
}
```

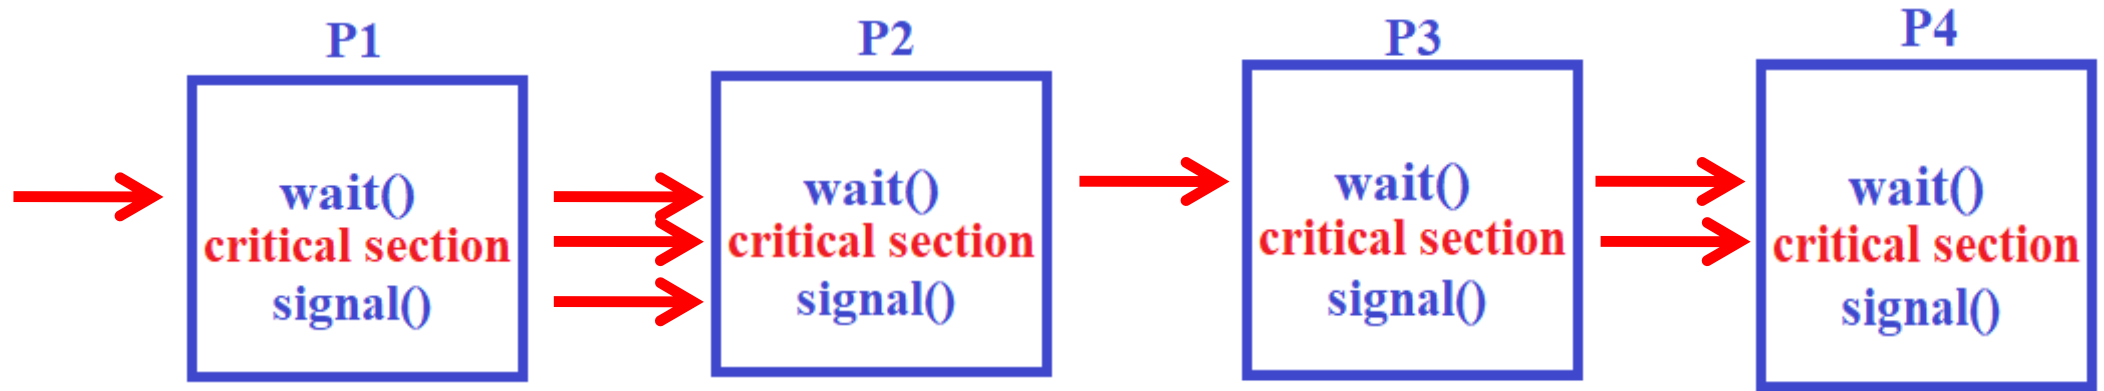
The definition of signal() is as follows:

```
signal(S) {  
    S++;  
}
```

- Each process that wishes **to use a resource** performs a **wait() operation** on the semaphore (thereby decrementing the count).
- When a process **leaves critical section**, it performs a **signal()** operation (incrementing the count).
- As long as the count for the semaphore goes to ≥ 0 , process invoking wait is allowed to enter critical section
- After count becomes < 0 , processes that wish to use a resource will busy wait until the count becomes greater than 0.

semaphore value

0



Blocking instead of busy waiting

- Recall that the implementations of **mutex**, **peter**, **compare & swap**, **test & test** are all suffering from **busy waiting**.
- The definitions of the **wait()** and **signal()** semaphore operations just described present **the same problem**.
- To overcome the need for busy waiting, we can modify the definition of the **wait()** and **signal()** operations as follows:
 - When a process **executes the wait()** operation and finds that the semaphore value is not positive, it must wait.
 - However, rather than engaging in busy waiting, the process can **block itself**.

- The block operation places a process into a **waiting queue** associated with the semaphore, and the state of the process is switched to the **waiting state**.
- Then control is transferred to the **CPU scheduler**, which **selects another process** to execute – **context switch**.

- A process that is **blocked**, waiting on a semaphore S, should be **restarted** when some other process executes a **signal()** operation.
- The process is restarted by a **wakeup()** operation, which changes the process from the **waiting state** to the **ready state**.
- The process is then placed in the **ready queue**.
- To implement semaphores under this definition, we define a semaphore as follows:

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore
```

- Now, the **wait()** semaphore operation can be defined as:

```
wait(semaphore *S) {
```

```
    S->value--;
```

```
    if (S->value < 0) { add this process to S->list;
```

```
        block();
```

```
    }
```

```
}
```

- and the **signal()** semaphore operation can be defined as

```
signal(semaphore *S) {
```

```
    S->value++;
```

```
    if (S->value <= 0) {
```

```
        remove a process P from S->list;
```

```
        wakeup(P);
```

```
    } }
```

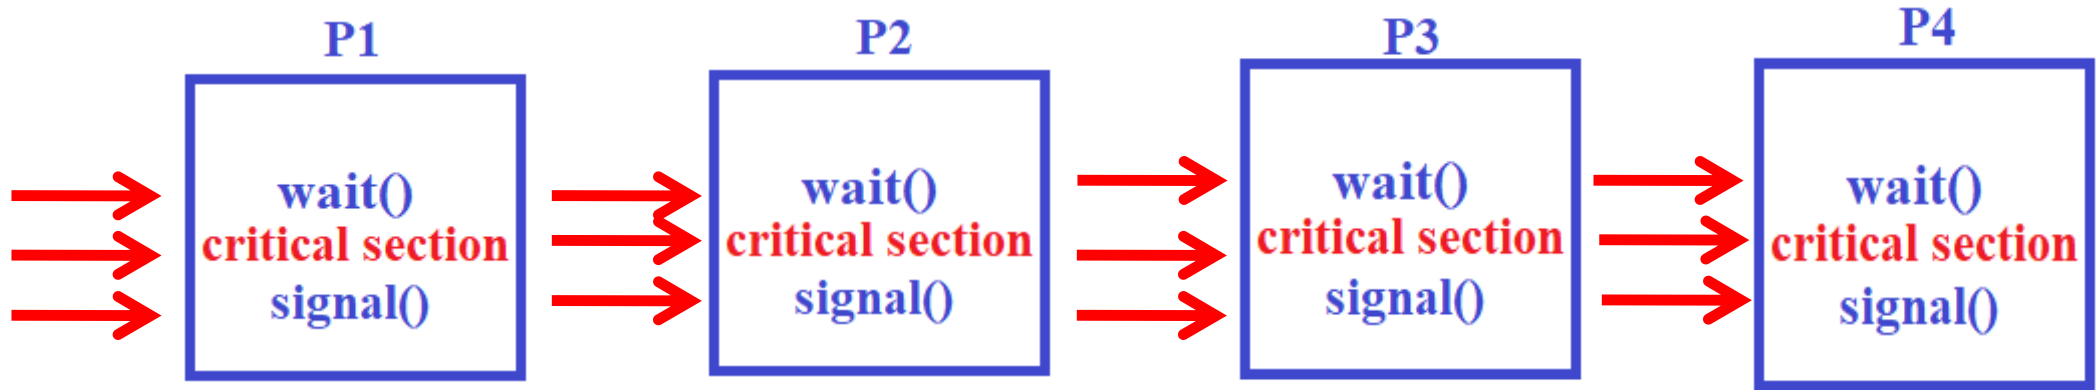

- The **block()** operation **suspends the process** that invokes it.
- The **wakeup(P)** operation **resumes the execution** of a blocked process P.
- These two operations are provided by the operating system as **basic system calls**

- In this implementation, semaphore values **may be negative**.
- If a semaphore value is **negative**, its **magnitude** is the **number of processes waiting** on that semaphore.
- The list of waiting processes can be easily implemented by a link field in each **process control block (PCB)**.
- Each semaphore contains an integer value and a pointer to a list of PCBs.
- One way to add and remove processes from the list so as to ensure bounded waiting is to use a **FIFO queue**, where the semaphore contains both head and tail pointers to the queue.
- In general, however, the list can use any queueing strategy

Blocked Queue





semaphore value



Sequencing instructions using Semaphores

- We can also use semaphores **to solve various other synchronization** problems.
- For example, consider two concurrently running processes: ***P1 with a statement S1 and P2 with a statement S2.***
- ***Suppose we require that S2 be executed only after S1 has completed.***
- ***We can implement this scheme readily by letting P1 and P2 share a common semaphore **synch**, initialized to 0.***

**P2 should execute s2 only after P1
completes s2**



synch = 0

P1

```
s1;  
signal(synch)  
s2;
```

P2

```
s1;  
wait(synch)  
s2;
```


- Initialize *synch* = 0

- *In process P1*, we insert the statements:

S1;

signal(synch);

- *In process P2*, we insert the statements

wait(synch);

S2;

- Because *synch* is initialized to 0, *P2* will execute *S2* only after *P1* has invoked *signal(synch)*, which is after statement *S1* has been executed.

Semaphore Types

- Operating systems often distinguish between counting and binary semaphores
- The value of a **counting semaphore** can range over an unrestricted domain
- The value of a **binary semaphore** can range only between 0 and 1

Advantage of Counting semaphore

- Counting semaphores can be used to control access to a given resource consisting of a **finite number of instances**.
- The semaphore is **initialized to the number of resources** available.
- Each process that wishes to use a resource **performs a wait() operation** on the semaphore (thereby decrementing the count).
- When a process releases a resource, it **performs a signal()** operation (incrementing the count).
- When the count for the **semaphore goes to 0**, all resources are being used.

- After that, processes that wish to use a resource **will block** until the count becomes greater than 0.
- Whereas **binary semaphores** can only be used when the **number of instances of the resource is exactly 1**.
- With counting semaphore, we may exactly **know number of processes waiting** for the critical section.

Ensuring Atomicity

- It is critical that semaphore operations be executed atomically.
- We must guarantee that **no two processes can execute wait() and signal() operations** on the same semaphore at the same time.
- This is a critical-section problem; and in a **single-processor** environment, we can solve it by simply **disabling interrupts** during the time the wait() and signal() operations are executing.
- This scheme **works in a single-processor** environment because, once interrupts are inhibited, instructions from **different processes cannot be interleaved**.
- Only the currently running process executes until interrupts are **re-enabled** and the scheduler can regain control.

- In a **multiprocessor environment**, interrupts must be **disabled on every processor**.
- Disabling interrupts on every processor can be a **difficult task** and furthermore can seriously **diminish performance**.
- Therefore, SMP systems must provide alternative locking techniques

Deadlocks and Starvation

- The implementation of a semaphore with a waiting queue may result in a situation where two or more processes are waiting indefinitely for an event that can be caused only by one of the waiting processes.
- When such a state is reached, these processes are said to be **deadlocked**.
- We say that a set of processes is in a deadlocked state when every process in the set is waiting for an event that can be caused only by another process in the set.

P_0

wait(S);

wait(Q);

.

.

.

signal(S);

signal(Q);

P_1

wait(Q);

wait(S);

.

.

.

signal(Q);


signal(S);

- Another problem related to deadlocks is **indefinite blocking or starvation**, a situation in which processes wait indefinitely within the semaphore.
- Indefinite blocking may occur if we remove processes from the list associated with a semaphore in LIFO (last-in, first-out) order.

Semaphore Primitives

```
struct semaphore {
    int count;
    queueType queue;
};
void semWait(semaphore s)
{
    s.count--;
    → if (s.count < 0) {
        /* place this process in s.queue */;
        /* block this process */;
    }
}
void semSignal(semaphore s)
{
    s.count++;
    → if (s.count <= 0) {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}
```

s.count



P1

**SemWait()
Critical_Section
SemSignal()**

P2

**SemWait()
Critical_Section
SemSignal**

Figure 5.3 A Definition of Semaphore Primitives

Binary Semaphore Primitives

```
struct binary_semaphore {
    enum {zero, one} value;
    queueType queue;
};
void semWaitB(binary_semaphore s)
{
    if (s.value == one)
        s.value = zero;
    else {
        /* place this process in s.queue */;
        /* block this process */;
    }
}
void semSignalB(semaphore s)
{
    if (s.queue is empty())
        s.value = one;
    else {
        /* remove a process P from s.queue */;
        /* place process P on ready list */;
    }
}
```

Figure 5.4 A Definition of Binary Semaphore Primitives

Priority Inversion

- A scheduling challenge arises when a **higher-priority process** needs to **access a resource** that is currently being accessed **by a lower-priority process**
- Since **shared resources** are typically protected with a **lock**, the **higher-priority process** will have to **wait** for a **lower-priority** one to finish with the resource
- The situation becomes more complicated **if the lower-priority process is preempted** in favor of another process with a higher priority.

- As an example, assume we have three processes—L, M, and H—whose priorities follow the **order $L < M < H$** .
- Assume that process **H requires resource R**, which is currently being **accessed by process L**.
- Ordinarily, **process H would wait for L** to finish using resource R.
- However, now suppose that process **M becomes runnable**, thereby **preempting process L**.
- Indirectly, a process with a **lower priority—process M—has affected** how long process **H must wait** for L to relinquish resource R.
- This problem is known as **priority inversion**

A high priority Ambulance vehicle is forced to wait behind the low priority auto-rickshaws due to traffic



Priority-inheritance protocol

- Typically these systems solve the problem by implementing a **priority-inheritance protocol**.
- According to this protocol, all processes that are accessing resources needed by a higher-priority process **inherit the higher priority** until they are finished with the resources in question.
- When they are finished, their **priorities revert** to their original values
- In the example above, a priority-inheritance protocol would allow process *L* to temporarily inherit the priority of process *H*, thereby preventing process *M* from preempting its execution

Classical Problems of Synchronization

- Classical problems that involve critical section problem:
 - Bounded-Buffer Problem
 - Readers and Writers Problem
 - Dining-Philosophers Problem

Producer - Consumer Problem

- **General Situation:**
 - One or more producers are **generating data** and **placing these** in a **buffer**
 - A **single consumer** is **taking items** out of the buffer one at time
 - **Only one producer or consumer** may access the buffer at any one time.
- **Conditions:**
 - Either the producer or the consumer should access the buffer at a time - ***Mutual exclusion***
 - Ensure that the ***Producer*** can't add data into ***full buffer*** and ***consumer*** can't remove data from ***empty buffer***.

Two variants of the P-C problem

- **Unbounded buffer**
 - Required conditions – ME, Empty buffer
- **Bounded buffer**
 - Required conditions – ME, Empty buffer, Full buffer

Solution for Bounded buffer problem

- In our problem, the producer and consumer processes share the following data structures:
- `int n;`
- `semaphore mutex = 1;` // to enforce mutual exclusion
- `semaphore empty = n;` // to prevent producer when buffer full
- `semaphore full = 0 ;` // to prevent consumer when buffer empty

mutex 1

Consumer

empty 5 full -1

```
do{  
→ wait(full) ;  
  wait(mutex);  
  .....  
  /*remove an item from buffer to next_consumed */  
  .....  
  signal(mutex);  
  signal(empty);  
  /*consume the item in next_consumed */  
}while(true);
```

Consumer blocked on queue
of Full semaphore

mutex 1

Producer

empty 8 full -1

```
do{  
→ wait(empty);  
→ wait(mutex);  
.....  
→ /* add an item into the buffer */  
.....  
→ signal(mutex);  
→ signal(full );  
}while(true);
```

Consumer unblocked from
of Full semaphore queue

mutex 0

Consumer

empty 5 full 0

```
do{  
    wait(full) ;  
→ wait(mutex);  
    .....  
→ /*remove an item from buffer to next_consumed */  
    .....  
→ signal(mutex);  
→ signal(empty);  
    /*consume the item in next_consumed */  
  
}while(true);
```

The Readers–Writers Problem

- Suppose that a **database** or **file** is to be **shared** among several concurrent processes.
- Some of these processes may **want only to read** the database, whereas others may **want to update** (that is, to read and write) the database.
- We distinguish between these two types of processes by referring to the former as ***readers*** and to the latter as ***writers***.
- if **two readers access** the shared data simultaneously, **no adverse effects** will result.
- However, **if a writer** and some other process (either a reader or a writer) access the database simultaneously, **chaos may ensue**.

- To ensure that these difficulties do not arise, we require that the **writers have exclusive access** to the shared database while writing to the database.
 - Mutual exclusion for writer
 - No mutual exclusion among readers
- This synchronization problem is referred to as the **readers–writers problem**.

Producer & Consumer Vs Reader & Writer

P- C	R-W
The buffer is filled by producer and it is emptied by consumer	Reader simply reads the values of a file and the writer updates the values
Mutual exclusion needed among any two processes whether two producers or two consumers or a producer and a consumer	No mutual exclusion required among readers since they only read, whereas writer requires exclusive access.
Two variants – bounded and unbounded buffer	Two variants – reader priority and writer priority

- **Variants**

- ***first readers–writers problem***, requires that no reader be kept waiting unless a writer has already obtained permission to use the shared object.
 - **second readers –writers problem** requires that, once a writer is ready, that writer perform its write as soon as possible.
- A solution to either problem may result in **starvation**
- In the first case, **writers may starve**; in the second case, **readers may starve**.
- For this reason, other variants of the problem have been proposed

- In the solution to the first readers–writers problem, the reader processes share the following data structures:
- **semaphore rw_mutex = 1;**
- **semaphore mutex = 1;**
- **int read count = 0;**
- The semaphore **rw mutex** is common to both reader and writer processes
- The **mutex** semaphore is used to ensure mutual exclusion when the variable read count is updated
- The **read_count** variable keeps track of how many processes are currently reading the object.

Readers have Priority

READER

```
do{
    wait(mutex);
    read_count++;
    if(read_count==1)
        wait(rw_mutex);
    signal(mutex)
    .....
    /*reading is performed */
    .....
    wait(mutex);
    read_count--;
    if(read_count==0)
        signal(mutex);
}while(true);
```

mutex – to prevent multiple readers accessing read_count variable
rw_mutex – to enforce ME among readers and writer

WRITER

```
do{
    wait(rw_mutex);
    .....
    /*writting is performed */
    .....
    signal(rw_mutex);
}while(true)
```

mutex

0

readcount

0

rw_mutex

-1

READER

```
do{  
    → wait(mutex);  
    → read_count++;  
    → if(read_count==1)  
        → wait(rw_mutex);  
    → signal(mutex)  
    .....  
    → /*reading is performed */  
    .....  
    → wait(mutex);  
    → read_count--;  
    → signal(mutex)  
    → if(read_count==0) ;  
        → signal(rw_mutex)  
}while(true);
```

Reader 1 -> CS

Reader 2 -> CS

Writer 1 -> CS

WRITER

```
do{  
    → wait(rw_mutex);  
    .....  
    → /*writting is performed */  
    .....  
    → signal(rw_mutex);  
  
}while(true)
```

Blocked Queue

W1				
----	--	--	--	--